

Question 1: Edge Detection vs. Harris Corner Detection for Satellite Image Registration

1.1 Case Study Context

A remote sensing application aligns satellite images of the same geographical region captured at different times, so that environmental changes (deforestation, urban growth, flooding, glacier retreat, etc.) can be tracked. Registration success depends entirely on finding a set of matching reference points between the two images that describe the same physical location on the ground. The system originally used edge detection as this reference-point source; in practice this proved unreliable because of illumination differences, orientation changes between passes, and cluttered background detail.

1.2 Why Edge Detection Fails in This Application

Edges mark locations where pixel intensity changes sharply along one direction, typically at object boundaries such as roads, field borders, coastlines, or building outlines. While edges are cheap to compute, they carry three structural weaknesses that make them a poor basis for image registration:

(a) The Aperture Problem — Ambiguous Localization Along the Edge

Within a small observation window, a straight edge looks identical at every point along its length. A matching algorithm therefore cannot tell which specific point on one image's edge corresponds to which point on the other image's edge — it can only tell that both windows contain 'an edge going this way.' This is illustrated below: a flat region has no gradient at all, an edge region has a gradient in only one direction (so motion along the edge is invisible to the detector), and only a genuine corner has gradients in two independent directions, which is what makes a point uniquely locatable.




Figure 1. The aperture problem: a flat region and an edge region cannot be localized uniquely inside a small window, while a corner-like structure can.

(b) Sensitivity to Illumination and Seasonal Differences

Satellite images of the same region taken at different times almost always differ in sun angle, shadow length, cloud cover, and seasonal ground cover (e.g., bare soil vs. crops). Edge detectors such as Sobel or Canny respond to intensity gradients, so a boundary that produces a strong edge in one image may weaken, shift, or disappear entirely in the other image purely because of lighting, even though the underlying terrain has not changed. This directly corrupts the correspondence step of registration.

(c) Repetitive and Cluttered Edge Structures

Agricultural boundaries, road grids, and coastlines generate many long, parallel, or repeating edges. Because of the aperture problem described above, these repetitive structures are easily confused with one another during matching, especially once the second image is slightly rotated, scaled, or translated relative to the first, as is normal between satellite passes.




Figure 2. Two edge maps of the same scene captured at different times: the edges themselves are similar, but their sheer repetitiveness and the shift/rotation between passes make point-to-point correspondence unreliable.

(d) No Natural Descriptor for Matching

An edge pixel by itself (position + gradient direction) is a weak descriptor — many other edge pixels in the same image look the same. This produces a large number of false candidate matches, which the registration algorithm must then filter out, increasing computational cost and reducing final alignment accuracy.

1.3 How Harris Corner Detection Improves the System

The Harris Corner Detector locates points where the local intensity surface changes significantly in two independent directions simultaneously (i.e., true corners: intersections of field boundaries, building corners, road junctions, distinctive terrain breaks). It does this by analyzing, within a small window, the structure tensor (second-moment matrix) of image gradients:

$$\mathbf{M} = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix}$$

The eigenvalues (λ_1, λ_2) of $\mathbf{M}$ describe how the intensity changes around a point in the two principal directions. Harris avoids the expensive eigenvalue computation by using a cornerness score:

$$\mathbf{R} = \det(\mathbf{M}) - k \cdot \text{trace}(\mathbf{M})^2 \quad (\text{typically } k \approx 0.04\text{--}0.06)$$

A large positive $\mathbf{R}$ indicates a corner (both λ_1 and λ_2 are large); a negative $\mathbf{R}$ indicates an edge (one eigenvalue much larger than the other); and a small $\mathbf{R}$ (both eigenvalues small) indicates a flat region. This is shown below for a simulated corner: the response peaks sharply and only at the true corner point, is strongly negative along the two edges, and is near zero in the flat background.




Figure 3. Simulated Harris cornerness response $R(x, y)$. The score is high only at the true corner location, confirming precise localization in two dimensions.

Specific Improvements Over Edge Detection

- Precise, two-dimensional localization: because a corner requires gradient change in two directions, its position is fixed exactly — there is no aperture-problem ambiguity, so the same physical point can be re-identified confidently in both images.
- Rotation invariance: the eigenvalues of the structure tensor (and hence the R score) do not change when the local patch is rotated, so corners detected in an image taken at one orientation are still detected — at the same physical locations — when the second pass arrives at a slightly different satellite orientation.
- Greater robustness to illumination change: Harris relies on relative gradient structure rather than absolute intensity, and using a Gaussian-weighted window further suppresses response to uniform brightness shifts and mild shadowing, so the same ground features are detected under different lighting.
- Distinctiveness for matching: corners (e.g., a road intersection, the corner of a building, the junction of two field boundaries) are far less repetitive across an image than generic edge pixels, which sharply reduces false-match candidates and improves the accuracy of the correspondence stage.
- Reduced computational load downstream: because Harris returns a sparse set of high-confidence keypoints instead of dense edge maps, subsequent descriptor computation, matching, and transform (homography) estimation all operate on far fewer, more reliable points.

Within the full registration pipeline, Harris corners typically feed into a descriptor stage (or are combined with a scale/rotation-invariant descriptor) before matching and geometric transform estimation, as shown below:




Figure 4. Image registration pipeline using Harris Corner Detection as the feature-extraction stage.

1.4 Comparative Summary

Criterion	Edge Detection	Harris Corner Detection
Localization	Ambiguous along the edge direction (aperture problem)	Precise 2-D point localization
Illumination changes	High sensitivity — gradients shift with lighting/shadows	More robust — relies on local gradient structure, not absolute intensity
Orientation changes	Edge maps can misalign after rotation between passes	Rotation-invariant cornerness response
Distinctiveness for matching	Low — many similar edge pixels cause false matches	High — corners are sparse and locally unique
Density of features	Dense (whole boundaries)	Sparse, high-confidence keypoints
Downstream computational cost	Higher (more candidate points to filter)	Lower (fewer, stronger points to match)
Suitability for registration	Poor under varying conditions	Well suited — improves accuracy and robustness

1.5 Conclusion for Question 1

Edge detection is limited in this application because it identifies boundaries rather than uniquely locatable points; it suffers from the aperture problem, is sensitive to illumination and seasonal changes between satellite passes, and produces repetitive, non-distinctive features that are hard to match reliably. Harris Corner Detection resolves these issues by isolating points with two-dimensional gradient variation, giving precise localization, strong invariance to rotation, useful robustness to illumination changes, and highly distinctive keypoints for matching. As a result, replacing edge detection with Harris Corner Detection directly improves feature localization accuracy, matching reliability, and the overall robustness of image registration under the varying imaging conditions typical of multi-temporal satellite data.

Question 2: Choosing an Image Representation for Real-Time Object Detection

2.1 Case Study Context

An autonomous surveillance system must detect vehicles and pedestrians in real time from high-resolution color video. Processing full-color frames directly is computationally expensive and slows down feature extraction, threatening the system's ability to meet real-time processing constraints. The engineering team is considering converting frames to grayscale or binary representations before running feature detection, to reduce computational load.

2.2 Comparing Image Representations

The three candidate representations differ in how much information they retain per pixel, and therefore in both computational cost and the quality of features that can be extracted from them.




Figure 5. The same simulated traffic scene represented as color (3 channels), grayscale (1 channel), and binary (1 bit per pixel).

(a) Color (RGB) Representation

- Retains full chromatic information (3 channels $\times$ 8 bits typically), which is essential for tasks such as distinguishing brake-light colors, traffic-signal states, or vehicle color/classification.
- Computationally the most expensive: convolution, gradient, and feature-detection operations must be run on three channels (or a merged higher-dimensional feature space), roughly tripling raw pixel throughput requirements compared with grayscale.
- More sensitive to illumination and color-balance variation (e.g., sodium-vapor street lighting, headlight glare) unless careful color-space normalization is applied.

(b) Grayscale Representation

- Reduces the image to a single intensity channel, typically via a weighted sum (e.g., $0.299R + 0.587G + 0.114B$), cutting per-pixel data volume roughly to one-third of color.
- Preserves the intensity gradients and structural/textural information that edge detectors (Sobel, Canny) and corner/feature detectors (Harris, FAST, HOG) actually operate on — most classical detectors are defined on single-channel intensity data in the first place.
- Substantially faster than color processing while retaining almost all of the shape and edge information needed to detect object boundaries (vehicle outlines, pedestrian silhouettes).
- Loses color-based cues (e.g., cannot distinguish a red car from a green car of the same shape), which is usually acceptable for detection/localization tasks but not for color-dependent classification.

(c) Binary Representation

- Reduces each pixel to 1 bit via thresholding, giving the lowest possible computational and memory cost and enabling very fast operations such as connected-component analysis and contour extraction.
- Extremely sensitive to the choice of threshold and to lighting conditions — a shadow, headlight glare, or overcast sky can shift the optimal threshold, causing objects to fragment or merge with the background.
- Discards fine gradient and texture information entirely, which degrades the performance of gradient-based feature detectors (Harris, Sobel/Canny, HOG) that rely on continuous intensity variation, not just an on/off silhouette.
- Best suited to controlled, high-contrast scenarios (e.g., fixed backlit conveyor inspection) rather than outdoor real-time traffic scenes with variable lighting.

2.3 Relative Computational Efficiency

Because color carries three channels and binary carries a single bit, the three representations differ substantially in raw data volume and downstream processing time. Using grayscale as the reference baseline, the approximate relative costs are illustrated below.




Figure 6. Approximate relative computational cost across the pipeline stages for each representation (illustrative, not measured benchmark data).

2.4 Comparative Summary

Criterion	Color (RGB)	Grayscale	Binary
Data per pixel	3 channels (~24 bits)	1 channel (~8 bits)	1 bit
Computational cost	Highest	Moderate	Lowest
Edge/gradient quality	Good, but redundant across channels	Good — matches classical detector assumptions	Poor — gradient information destroyed
Robustness to illumination	Sensitive without normalization	Reasonably robust	Highly sensitive to threshold/lighting
Feature extraction suitability (Harris/Sobel/HOG etc.)	Usable but costly; often converted internally to grayscale anyway	Well matched — most detectors are designed for this	Weak — fine structure lost
Color-dependent cues (signals, brake lights)	Preserved	Lost	Lost

Best fit for real-time detection	Only if color cues are essential and hardware allows	Best overall balance	Only for controlled/high-contrast setups
----------------------------------	--	----------------------	--

2.5 Justification and Recommendation

For a real-time vehicle and pedestrian detection system, grayscale representation is the most appropriate choice. The reasoning is as follows:

- Feature detectors used for object detection (edge detectors, Harris corners, HOG-style gradient descriptors) are fundamentally built on intensity gradients, not color; converting to grayscale removes redundant channel information without discarding the gradient structure these algorithms depend on.
- Grayscale cuts computational load roughly three-fold relative to full color, directly addressing the stated bottleneck in achieving real-time frame rates, while binary's further savings come at the cost of destroying the gradient information that feature detectors require, causing missed or fragmented detections under real-world lighting.
- Binary thresholding is only reliable in tightly controlled lighting; an outdoor, always-on surveillance system faces continuously varying illumination (day/night, headlights, shadows, weather), which makes a fixed or even adaptive threshold a fragile choice for a safety-relevant detection task.
- If color-dependent cues are genuinely required for part of the pipeline (e.g., traffic-light state or brake-light recognition), a hybrid approach can be used: run the primary, time-critical detection pipeline (edge/corner/HOG-based detection) on grayscale frames for speed, and apply targeted color analysis only within already-detected regions of interest, rather than across the full high-resolution color frame.

This hybrid, grayscale-first strategy achieves the best balance between computational efficiency, detection accuracy, and robustness, and is the representation strategy justified for this real-time computer vision application.

2.6 Conclusion for Question 2

Color, grayscale, and binary representations trade off information content against computational cost. Color preserves the most information but is the most expensive and least necessary for shape-based detection; binary is the cheapest but discards the gradient information that feature detectors rely on and is highly unstable under variable lighting. Grayscale offers the best overall balance — substantially reduced computational cost compared to color, while retaining the intensity-gradient structure required for reliable edge and feature-based detection — making it the most appropriate primary representation for accurate, real-time object detection, optionally supplemented with selective color analysis where color-specific cues are required.